

Enhancement & Customizations for

Workshop: RAG using Amazon Bedrock Agents and Knowledge Base

Description:

The original “**RAG using Amazon Bedrock Agents and Knowledge base**” workshop guides you how to create a GenAI bedrock agent that can answer questions about Amazon SageMaker.

This document will guide you how to enhance this workshop by:

- 1) Use your own data as the data source for the Bedrock agent.
- 2) Customize your own Agent & Knowledge Base instructions (Prompts)
- 3) Creating a webpage that will allow users from outside of this workshop environment to experiment with your deliverables and send prompts to the agent that you created.

Use your own data

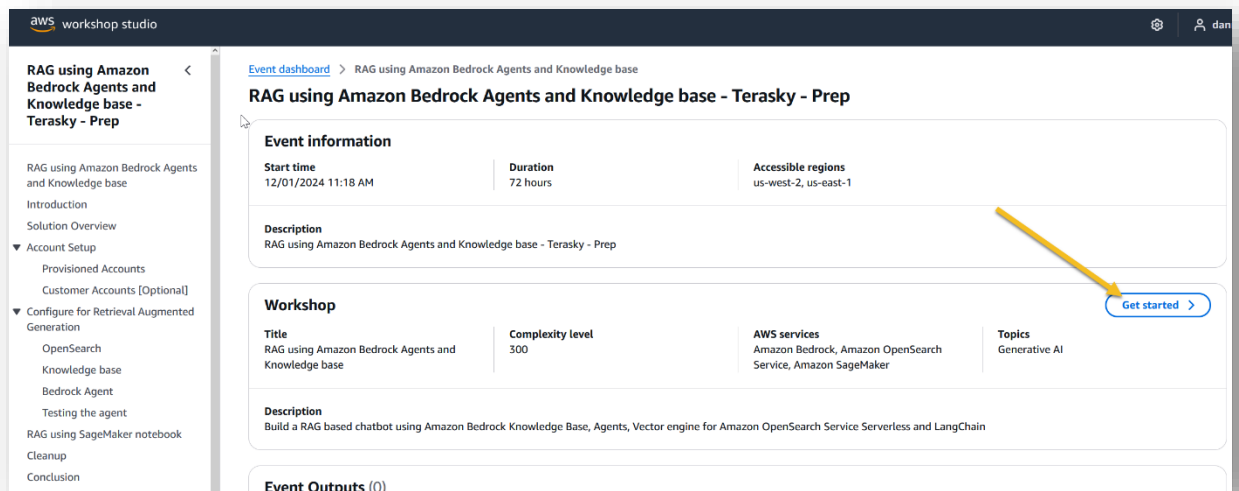
The workshop session starts with the deployment of the “**rag-w-bedrock-kb**” CloudFormation stack. (In AWS hosted events this stack is pre-deployed before you login to the workshop account)

The **rag-w-bedrock-kb** stack is used to define resources like... IAM roles, Access Policies, S3-Bucket and more... that you get to use in the workshop tasks.

The **rag-w-bedrock-kb** stack also import SageMaker-documentation HTML files into the S3-Bucket, and these files are used as the data source for creating a vector database that the bedrock agent will use.

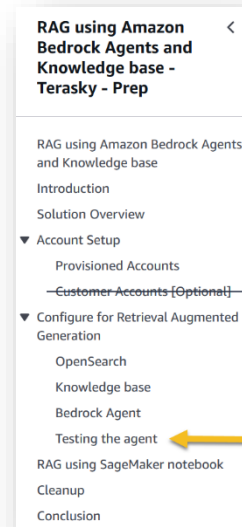
In order to use your own data in this workshop, we need to delete these SageMaker-documentation files and upload your files to the same S3-Bucket.

- 1) Once you have uploaded the relevant files to the S3-Bucket. Go to the Event dashboard page and click - **Get started** >



- 2) Follow the workshop lab instructions till you complete the **Testing the agent** phase.

Customer Accounts [Optional] phase is not required.



Important:

In these workshop steps you will be asked to define resources with specific names like “**sagemaker-qa**” (Bedrock agent) or “**sagemaker-docs**” (Bedrock Knowledge Base).

Please make sure that you use the exact names because the initial permissions rules that were defined in this environment are using these names.

Define your own agent & knowledge Base instructions

On step **Bedrock Agent** tasks 2 & 4 – you are asked to define **Instructions for the Agent & Knowledge base instructions for agent**

The suggested prompts are for a Q&A agent that answer to SageMaker related questions.

This is the time to get creative –

- What prompt will you use?
- What LLM works best for you?
- How can you make your prompt helpful for users?
- Will the prompt hold any output examples?

RAG using Amazon Bedrock Agents and Knowledge base - C-Level - Development

RAG using Amazon Bedrock Agents and Knowledge base

Introduction

Solution Overview

▼ Account Setup

- Provisioned Accounts
- Customer Accounts [Optional]

▼ Configure for Retrieval Augmented Generation

- OpenSearch
- Knowledge base
- Bedrock Agent**
- Testing the agent

Troubleshooting: If your prompt is not working like it should...

- Are you getting this info message?

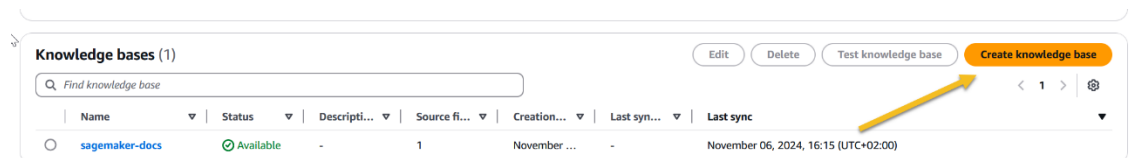
i To optimize agent performance, the agent doesn't use instructions if all of the following conditions are true:

- This agent has no action groups
- There is only one knowledge base associated to this agent
- No advanced prompts are overridden
- User input and code interpreter are disabled

If any of the above conditions is false, the agent will use the instructions. [Learn more](#)

In order to overcome this limitation let's add a second knowledge base to our agent.

- On the bedrock page, open the Knowledge base pane
- Click **Create Knowledge base**



- Name the knowledge base “knowledge-empty” and click **Next**

Provide knowledge base details

Knowledge base details

Knowledge base name

knowledge-empty

Valid characters are a-z, A-Z, 0-9, _ (underscore) and - (hyphen). The name can have up to 50 characters.

Knowledge base description - optional

Enter description

Valid characters are a-z, A-Z, 0-9, _ (underscore) and - (hyphen). The name can have up to 200 characters.

- Name the Data source “data-source-empty”, Select the S3-Bucket and click **Next**

Data source: data-source-empty Delete

Data source name
data-source-empty
Valid characters are a-z, A-Z, 0-9, _ (underscore) and - (hyphen). The name can have up to 100 characters.

Data source location
☒ This AWS account
☐ Other AWS account

S3 URI
To increase the accuracy and relevance of your responses, add a `meta_data.json` file containing metadata for your data source to your S3 bucket. [Info](#)

S3 URI
 View Browse S3

☐ Add customer-managed KMS key for S3 data - optional
If you encrypted your S3 data, provide the KMS key here so that Bedrock can decrypt it.

Chunking and parsing configurations [Info](#)
Choose between default or advanced customization.

☒ **Default**
Uses default parsing and chunking strategy.
 ☐ **Custom**
Customize the parsing and chunking strategy, including using advanced parsing.

► **Advanced settings - optional**

Add data source
You can add 4 more data source(s).

Cancel Previous Next

- Select the **Titan Text Embeddings V2** model and click **Next**.
(Without setting the vector DB)

Embeddings model
Select an embeddings model to convert your data into an embedding. Pricing depends on the model. [Learn more](#)

☒ **Titan Text Embeddings v2** [By Amazon](#)
☐ **Titan Embeddings G1 - Text v1.2** [By Amazon](#)
☐ **Embed English v3** [By Cohere](#)
☐ **Embed Multilingual v3** [By Cohere](#)

Vector dimensions
Select the vector dimension size for your embeddings model to balance accuracy, cost, and latency. Higher dimensions improves overall accuracy and requires more vector storage. [Learn more](#)

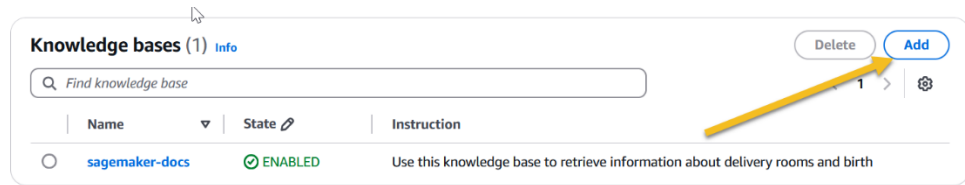
Delete Edit Create Agent

- Click **Create Knowledge base** and wait for the creation to be completed
- Navigate back to the **Bedrock** → **Agents** pane
- Select the **sagemaker-qa** agent and click **Edit**

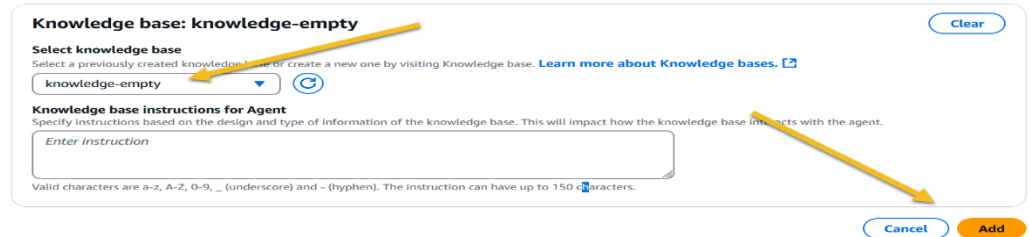
Agents (1) Delete Edit Create Agent

Name	Status	Description	Last updated
<input checked="" type="radio"/> sagemaker-qa	Prepared		November 07, 2024, 12:30 (UTC+02:00)

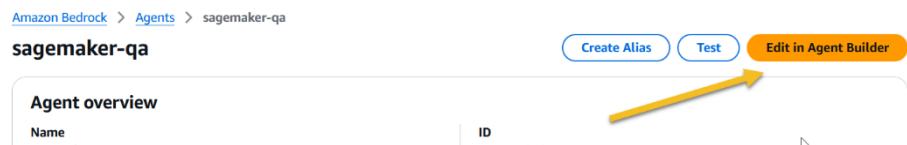
- Under the **Knowledge base** pane click **add**



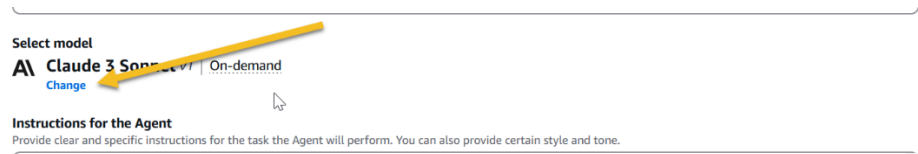
- Select the empty Knowledge base that you created and click **Add**



- Click **Save and Exit** in the Agent Builder page
- Click **Prepare**
- Test your agent...
- Still not getting the reply that you are expecting...
 - Click **Edit in Agent Builder**



- Click **change** and select a different model



- **Save** the agent, **Prepare** the agent and test it
- Experiment with different models till you get the required results

Prompt example:

- **Instructions for the Agent:**

You are a medical assistant by the name Daniela.

Your task is to answer patient's questions based on the data in the knowledge base only.

The answer should be written in simple language in an E-mail format with new lines between sentences to format a well-structured email.

The context of the questions is always around birth and delivery rooms.

Example email reply:

Dear patient,

It is recommended that the expectant mother eat snacks that she enjoys during childbirth, as the labor process can be long and the mother may feel hungry.

However, heavy meals are not recommended, especially if the mother is receiving an epidural, as she can only drink water during that time.

I recommend packing light snacks in the delivery room bag to help sustain the mother's energy during labor.

Thank you and see you at the delivery room,
Daniela

- **Knowledge base instructions for agent:**

Your task is to generate email based on the patient's query. Answers questions that are only based on the knowledge base and rephrase the answer in simple English.

Congratulations!

You have created a GenAI bedrock agent with an OpenSearch vector Database that hosts your embedded vectorized data as a RAG.

Now...

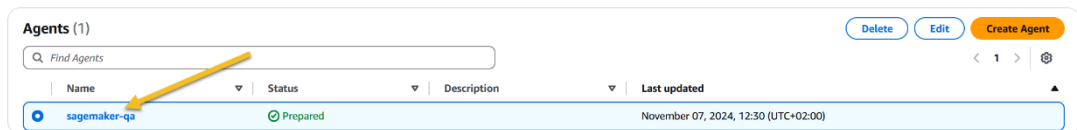
Let's expose your agent to the world

Expose your Agent to the world

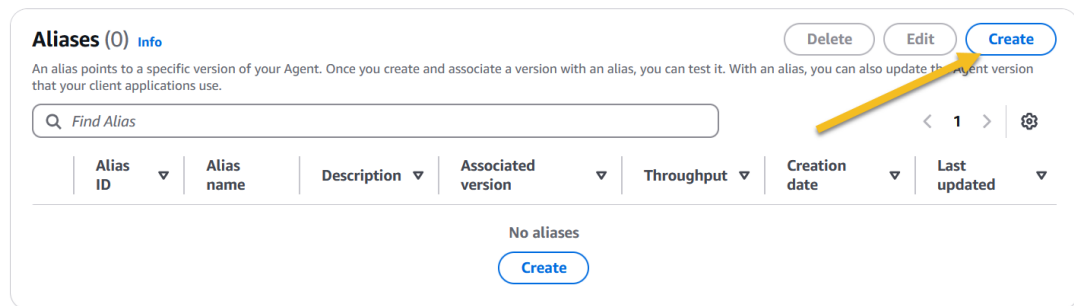
This section will instruct you on how to deploy a CloudFormation template that will create a webpage that will allow external users to send questions to & get-replies from your Bedrock Agent from outside the workshop environment.

In order to access our Bedrock agent with a lambda function, we have to define an Agent Alias.

- 1) Navigate to the **Bedrock** → **Agents** pane
- 2) Click on the **sagemaker-qa** agent



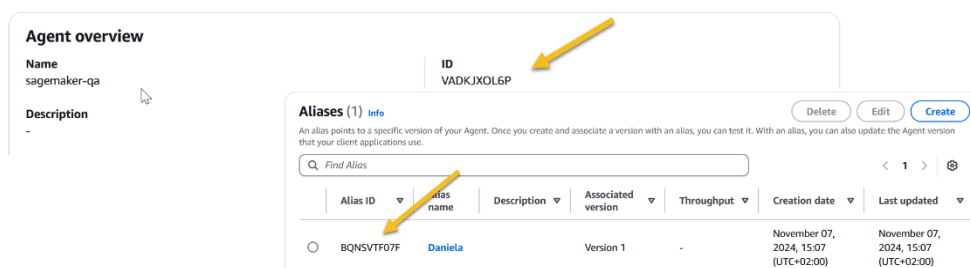
- 3) Scroll down to the Aliases pane and click **Create**



- 4) Type an **Alias name** and click **Create Alias**

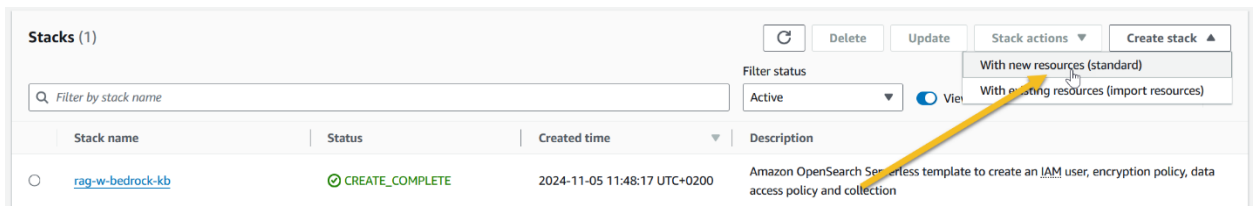
The screenshot shows the 'Create alias' dialog box. It has a text input field for 'Alias name' with the value 'Daniela'. Below it is a text area for 'Description - optional'. There are two radio buttons for 'Associate a version': 'Create a new version and associate it to this alias.' (selected) and 'Use an existing version to associate this alias.'. There are two radio buttons for 'Select throughput': 'On-demand' (selected) and 'Provisioned Throughput'. A yellow arrow points to the 'Alias name' field, and another yellow arrow points to the 'Create alias' button.

- 5) From the agent page copy the **Alias ID** & the **Agent ID**

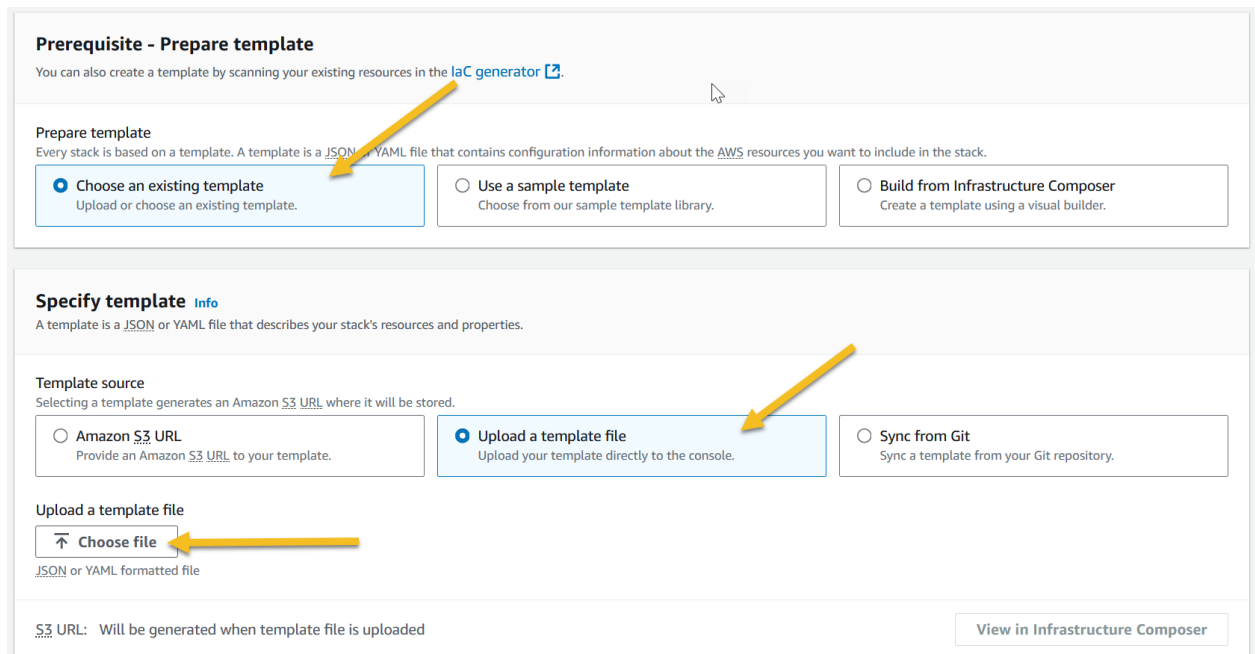


Deploy the Agent-Frontend CloudFormation

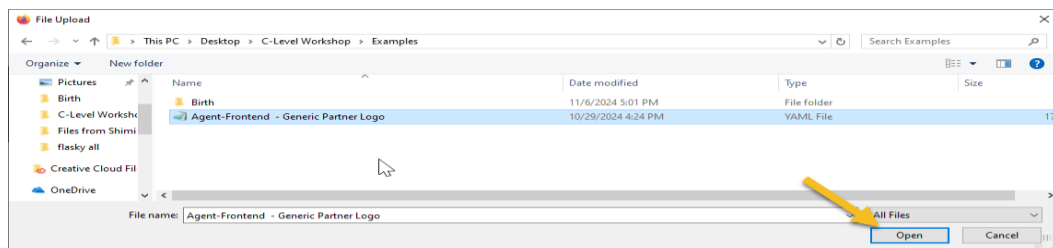
- 1) Navigate to the CloudFormation page and click **Create Stack** → **With new resources**



- 2) Select **Choose an existing template & upload a template file**, and click **Choose file**



- 3) Select the stack file that you got from the workshop facilitator and click open

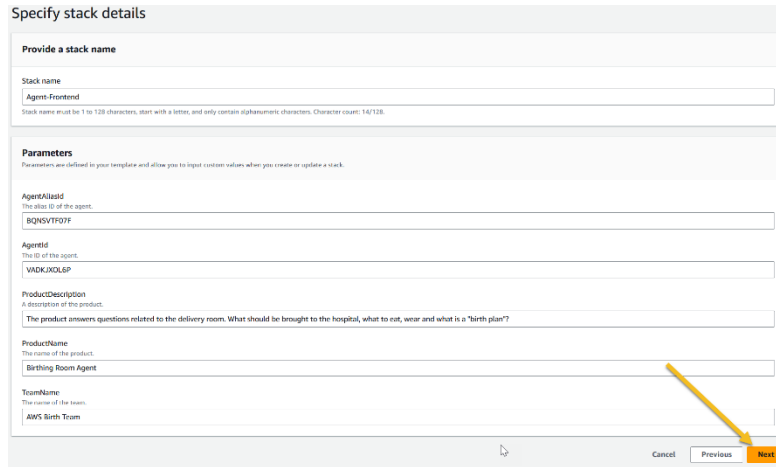


- 4) Click **Next**

5) Fill the required information

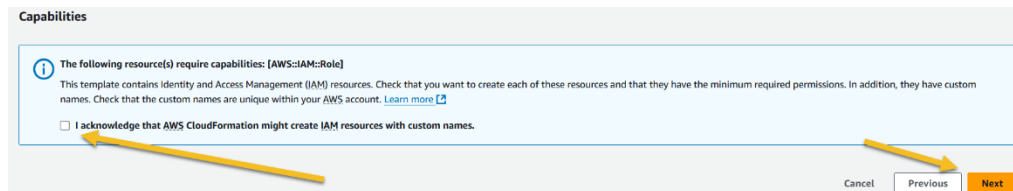
- a. **Stack name:** Agent-Frontend
- b. **AgentAliasId:** <The bedrock agent alias ID that we saved before>
- c. **AgentId:** <The bedrock agent ID that we saved before>
- d. **ProductDescription:** <A description of the product>
- e. **ProductName:** <The name of the product>
- f. **TeamName:** <Your Name/s>

And click **Next**



The image shows a 'Specify stack details' form in the AWS CloudFormation console. It has two main sections: 'Provide a stack name' and 'Parameters'. The 'Provide a stack name' section has a text input field for 'Stack name' with the value 'Agent-Frontend'. The 'Parameters' section has several input fields: 'AgentAliasId' (value: 'BQNSVTF07F'), 'AgentId' (value: 'VADKJOL6P'), 'ProductDescription' (value: 'The product answers questions related to the delivery room. What should be brought to the hospital, what to eat, wear and what is a "birth plan"?'), 'ProductName' (value: 'Birthing Room Agent'), and 'TeamName' (value: 'AWS Birth Team'). At the bottom right, there are 'Cancel', 'Previous', and 'Next' buttons. A yellow arrow points to the 'Next' button.

6) Acknowledge and click **Next**

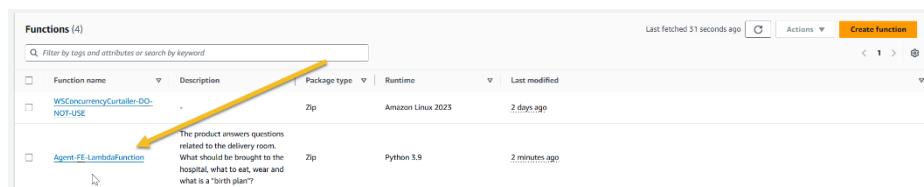


The image shows a 'Capabilities' screen in the AWS CloudFormation console. It contains a blue box with a warning icon and text: 'The following resource(s) require capabilities: [AWS::IAM::Role]. This template contains Identity and Access Management (IAM) resources. Check that you want to create each of these resources and that they have the minimum required permissions. In addition, they have custom names. Check that the custom names are unique within your AWS account. [Learn more](#)'. Below this, there is a checkbox labeled 'I acknowledge that AWS CloudFormation might create IAM resources with custom names.' which is currently unchecked. At the bottom right, there are 'Cancel', 'Previous', and 'Next' buttons. A yellow arrow points to the 'Next' button.

7) Review and click **Submit**

8) Wait 2-3 minutes till all the resources get created and navigate to the Lambda page

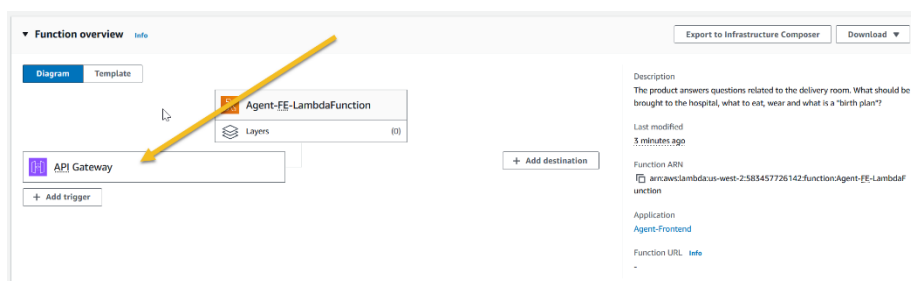
9) In the Lambda page, click on the **Agent-FE-LambdaFunction** function



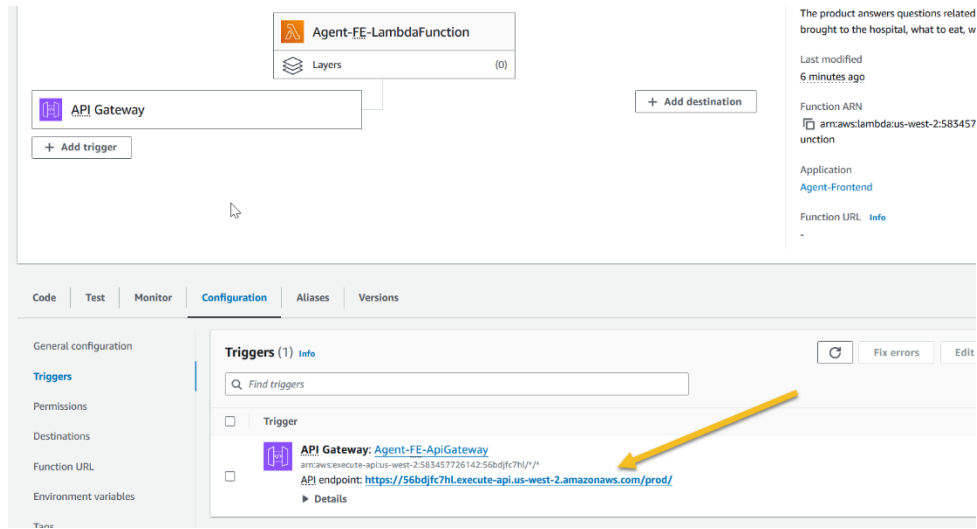
The image shows the 'Functions (4)' list in the AWS Lambda console. It has a search bar and a table with columns: Function name, Description, Package type, Runtime, and Last modified. The table contains four entries: 'WSConcurrencyCurtailed-DO-NOT-USE', 'Agent-FE-LambdaFunction', and two others. The 'Agent-FE-LambdaFunction' entry is highlighted with a yellow arrow. The description for 'Agent-FE-LambdaFunction' is 'The product answers questions related to the delivery room. What should be brought to the hospital, what to eat, wear and what is a "birth plan"?'. The package type is 'Zip', the runtime is 'Python 3.9', and it was last modified '2 minutes ago'.

Function name	Description	Package type	Runtime	Last modified
WSConcurrencyCurtailed-DO-NOT-USE		Zip	Amazon Linux 2023	2 days ago
Agent-FE-LambdaFunction	The product answers questions related to the delivery room. What should be brought to the hospital, what to eat, wear and what is a "birth plan"?	Zip	Python 3.9	2 minutes ago

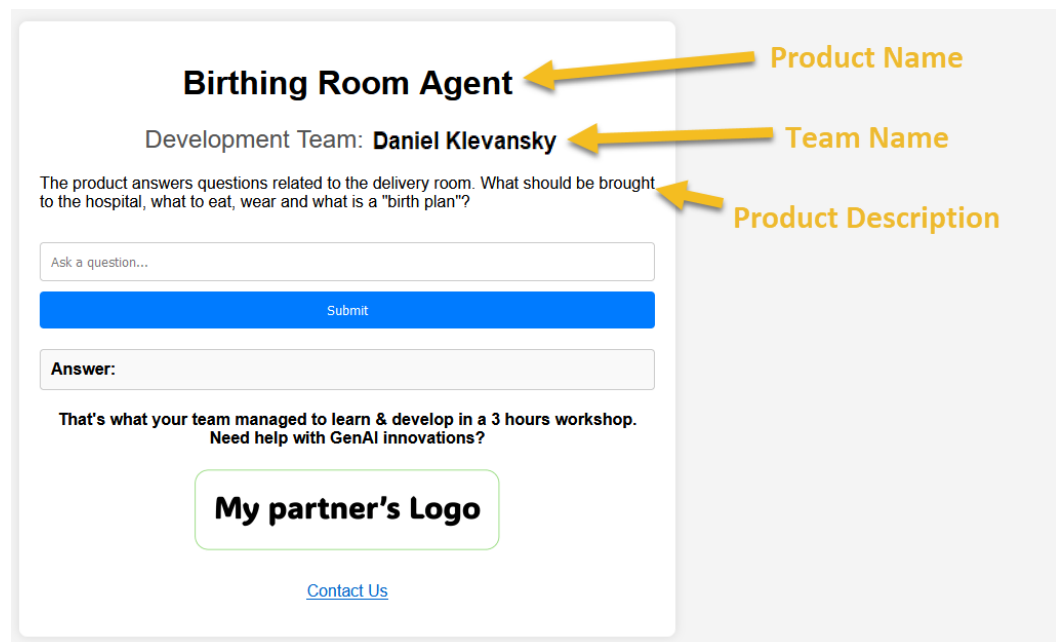
10) Click on the API Gateway component



11) Under configuration click on the API endpoint link and access your agent via the internet.



12) You should see a web page that looks like this...



13) Access this link from your phone & enjoy your first Bedrock Agent